

Data Preprocessing

4

8.7

10.6

| | Fuel Consumption Comb (mpg) | CO2 Emissions(g/km) |
|---|-----------------------------|---------------------|
| 0 | 33 | 196 |
| 1 | 29 | 221 |
| 2 | 48 | 136 |
| 3 | 25 | 255 |
| 4 | 27 | 244 |

df.head()

| | Make | Model | Vehicle Class | Engine Size(L) | Cylinders |
|--------------------|-------|------------|---------------|----------------|-----------|
| Transmission \ AS5 | ACURA | ILX | COMPACT | 2.0 | 4 |
| M6 | ACURA | ILX | COMPACT | 2.4 | 4 |
| AV7 | ACURA | ILX HYBRID | COMPACT | 1.5 | 4 |
| AS6 | ACURA | MDX 4WD | SUV - SMALL | 3.5 | 6 |
| AS6 | ACURA | RDX AWD | SUV - SMALL | 3.5 | 6 |

| | Fuel Type | Fuel Consumption City (L/100 km) \ |
|---|-----------|------------------------------------|
| 0 | Z | 9.9 |
| 1 | Z | 11.2 |
| 2 | Z | 6.0 |
| 3 | Z | 12.7 |
| 4 | Z | 12.1 |

| | Fuel Consumption Hwy (L/100 km) | Fuel Consumption Comb (L/100 km) \ |
|---|---------------------------------|------------------------------------|
| 0 | 6.7 | 8.5 |
| 1 | 7.7 | 9.6 |
| 2 | 5.8 | 5.9 |
| 3 | 9.1 | 11.1 |
| 4 | 8.7 | 10.6 |

| | Fuel Consumption Comb (mpg) | CO2 Emissions(g/km) |
|---|-----------------------------|---------------------|
| 0 | 33 | 196 |
| 1 | 29 | 221 |
| 2 | 48 | 136 |
| 3 | 25 | 255 |
| 4 | 27 | 244 |

```
df.shape
(7385, 12)
len(df)
7385
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Make              7385 non-null    object  
 1   Model             7385 non-null    object  
 2   Vehicle Class    7385 non-null    object  
 3   Engine Size(L)   7385 non-null    float64 
 4   Cylinders         7385 non-null    int64   
 5   Transmission      7385 non-null    object  
 6   Fuel Type          7385 non-null    object  
 7   Fuel Consumption City (L/100 km) 7385 non-null    float64 
 8   Fuel Consumption Hwy (L/100 km) 7385 non-null    float64 
 9   Fuel Consumption Comb (L/100 km) 7385 non-null    float64 
 10  Fuel Consumption Comb (mpg)       7385 non-null    int64   
 11  CO2 Emissions(g/km)            7385 non-null    int64  
dtypes: float64(4), int64(3), object(5)
memory usage: 692.5+ KB

df.isna().sum()

Make                  0
Model                 0
Vehicle Class         0
Engine Size(L)        0
Cylinders             0
Transmission          0
Fuel Type              0
Fuel Consumption City (L/100 km) 0
Fuel Consumption Hwy (L/100 km) 0
Fuel Consumption Comb (L/100 km) 0
Fuel Consumption Comb (mpg)       0
CO2 Emissions(g/km)      0
dtype: int64

df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384
Data columns (total 12 columns):
 #   Column           Non-Null Count Dtype  
 --- 
 0   Make             7385 non-null   object  
 1   Model            7385 non-null   object  
 2   Vehicle Class   7385 non-null   object  
 3   Engine Size(L)  7385 non-null   float64 
 4   Cylinders        7385 non-null   int64   
 5   Transmission     7385 non-null   object  
 6   Fuel Type        7385 non-null   object  
 7   Fuel Consumption City (L/100 km) 7385 non-null   float64 
 8   Fuel Consumption Hwy (L/100 km)   7385 non-null   float64 
 9   Fuel Consumption Comb (L/100 km) 7385 non-null   float64 
 10  Fuel Consumption Comb (mpg)      7385 non-null   int64  
 11  CO2 Emissions(g/km)            7385 non-null   int64  
dtypes: float64(4), int64(3), object(5)
memory usage: 692.5+ KB

```

```
df['Make'].value_counts()
```

| Make | |
|---------------|-----|
| FORD | 628 |
| CHEVROLET | 588 |
| BMW | 527 |
| MERCEDES-BENZ | 419 |
| PORSCHE | 376 |
| TOYOTA | 330 |
| GMC | 328 |
| AUDI | 286 |
| NISSAN | 259 |
| JEEP | 251 |
| DODGE | 246 |
| KIA | 231 |
| HONDA | 214 |
| HYUNDAI | 210 |
| MINI | 204 |
| VOLKSWAGEN | 197 |
| MAZDA | 180 |
| LEXUS | 178 |
| JAGUAR | 160 |
| CADILLAC | 158 |
| SUBARU | 140 |
| VOLVO | 124 |
| INFINITI | 108 |
| BUICK | 103 |
| RAM | 97 |
| LINCOLN | 96 |
| MITSUBISHI | 95 |

```
CHRYSLER      88
LAND ROVER    85
FIAT          73
ACURA         72
MASERATI      61
ROLLS-ROYCE   50
ASTON MARTIN 47
BENTLEY       46
LAMBORGHINI   41
ALFA ROMEO   30
GENESIS        25
SCION          22
SMART          7
BUGATTI        3
SRT            2
Name: count, dtype: int64
```

```
df[df['Make'] == 'SMART'][['Make', 'Fuel Type']]
```

```
   Make Fuel Type
943  SMART      Z
944  SMART      Z
2072 SMART      Z
2073 SMART      Z
3189 SMART      Z
3190 SMART      Z
4255 SMART      Z
```

```
df.columns
```

```
Index(['Make', 'Model', 'Vehicle Class', 'Engine Size(L)',  
'Cylinders',  
       'Transmission', 'Fuel Type', 'Fuel Consumption City (L/100  
km)',  
       'Fuel Consumption Hwy (L/100 km)', 'Fuel Consumption Comb  
(L/100 km)',  
       'Fuel Consumption Comb (mpg)', 'CO2 Emissions(g/km)'),  
      dtype='object')
```

```
df['Make'].value_counts()
```

```
Make
FORD      628
CHEVROLET 588
BMW       527
MERCEDES-BENZ 419
PORSCHE   376
```

```
TOYOTA      330
GMC         328
AUDI        286
NISSAN      259
JEEP         251
DODGE       246
KIA          231
HONDA        214
HYUNDAI     210
MINI         204
VOLKSWAGEN   197
MAZDA        180
LEXUS        178
JAGUAR       160
CADILLAC     158
SUBARU       140
VOLVO        124
INFINITI     108
BUICK        103
RAM           97
LINCOLN      96
MITSUBISHI   95
CHRYSLER     88
LAND ROVER    85
FIAT          73
ACURA         72
MASERATI      61
ROLLS-ROYCE   50
ASTON MARTIN  47
BENTLEY       46
LAMBORGHINI   41
ALFA ROMEO    30
GENESIS        25
SCION          22
SMART          7
BUGATTI         3
SRT            2
Name: count, dtype: int64
```

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df[ "Make" ] = le.fit_transform(df[ "Make" ])
df[ "Make" ].value_counts()
```

```
Make
13      628
9       588
5       527
28      419
32      376
```

```
39    330
15    328
3     286
31    259
20    251
11    246
21    231
16    214
17    210
29    204
40    197
27    180
24    178
19    160
8     158
38    140
41    124
18    108
7     103
33    97
25    96
30    95
10    88
23    85
12    73
0     72
26    61
34    50
2     47
4     46
22    41
1     30
14    25
35    22
36    7
6     3
37    2
Name: count, dtype: int64

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384
Data columns (total 12 columns):
 #   Column           Non-Null Count Dtype
 --- 
 0   Make            7385 non-null  int64
 1   Model           7385 non-null  object
 2   Vehicle Class  7385 non-null  object
 3   Engine Size(L) 7385 non-null  float64
```

```

4   Cylinders                      7385 non-null    int64
5   Transmission                   7385 non-null    object
6   Fuel Type                      7385 non-null    object
7   Fuel Consumption City (L/100 km) 7385 non-null    float64
8   Fuel Consumption Hwy (L/100 km) 7385 non-null    float64
9   Fuel Consumption Comb (L/100 km) 7385 non-null    float64
10  Fuel Consumption Comb (mpg)     7385 non-null    int64
11  CO2 Emissions(g/km)           7385 non-null    int64
dtypes: float64(4), int64(4), object(4)
memory usage: 692.5+ KB

df['Model'].value_counts()

Model
F-150 FFV                  32
F-150 FFV 4X4               32
MUSTANG                     27
FOCUS FFV                   24
SONIC                       20
..
Camry TRD                   1
Cullinan Black Badge        1
1500 4X4 EcoDiesel          1
1500 EcoDiesel               1
Sentra SR                    1
Name: count, Length: 2053, dtype: int64

df.drop('Model', axis =1, inplace=True)
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384
Data columns (total 11 columns):
 #   Column            Non-Null Count  Dtype  
 ---  -- 
 0   Make              7385 non-null    int64  
 1   Vehicle Class     7385 non-null    object  
 2   Engine Size(L)   7385 non-null    float64 
 3   Cylinders         7385 non-null    int64  
 4   Transmission      7385 non-null    object  
 5   Fuel Type         7385 non-null    object  
 6   Fuel Consumption City (L/100 km) 7385 non-null    float64 
 7   Fuel Consumption Hwy (L/100 km) 7385 non-null    float64 
 8   Fuel Consumption Comb (L/100 km) 7385 non-null    float64 
 9   Fuel Consumption Comb (mpg)     7385 non-null    int64  
 10  CO2 Emissions(g/km)           7385 non-null    int64  
dtypes: float64(4), int64(4), object(3)
memory usage: 634.8+ KB

df['Vehicle Class'].value_counts()

```

```

Vehicle Class
SUV - SMALL           1217
MID-SIZE              1133
COMPACT                1022
SUV - STANDARD         735
FULL-SIZE              639
SUBCOMPACT              606
PICKUP TRUCK - STANDARD 538
TWO-SEATER              460
MINICOMPACT              326
STATION WAGON - SMALL      252
PICKUP TRUCK - SMALL      159
MINIVAN                  80
SPECIAL PURPOSE VEHICLE    77
VAN - PASSENGER            66
STATION WAGON - MID-SIZE      53
VAN - CARGO                  22
Name: count, dtype: int64

df["Vehicle Class"] = le.fit_transform(df["Vehicle Class"])
df["Vehicle Class"].value_counts()

Vehicle Class
11    1217
2     1133
0     1022
12    735
1     639
10    606
6     538
13    460
3     326
9     252
5     159
4     80
7     77
15    66
8     53
14    22
Name: count, dtype: int64

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384
Data columns (total 11 columns):
 #   Column          Non-Null Count  Dtype  
 --- 
 0   Make             7385 non-null   int64  
 1   Vehicle Class   7385 non-null   int64  

```

```
2 Engine Size(L) 7385 non-null float64
3 Cylinders 7385 non-null int64
4 Transmission 7385 non-null object
5 Fuel Type 7385 non-null object
6 Fuel Consumption City (L/100 km) 7385 non-null float64
7 Fuel Consumption Hwy (L/100 km) 7385 non-null float64
8 Fuel Consumption Comb (L/100 km) 7385 non-null float64
9 Fuel Consumption Comb (mpg) 7385 non-null int64
10 CO2 Emissions(g/km) 7385 non-null int64
dtypes: float64(4), int64(5), object(2)
memory usage: 634.8+ KB
```

```
df[ 'Transmission' ].value_counts()
```

```
Transmission
```

```
AS6    1324
AS8    1211
M6     901
A6     789
A8     490
AM7    445
A9     339
AS7    319
AV     295
M5     193
AS10   168
AM6    132
AV7    118
AV6    113
M7     91
A5     84
AS9    77
A4     65
AM8    62
A7     53
AV8    39
A10   31
AS5    26
AV10   11
AM5    4
AM9    3
AS4    2
```

```
Name: count, dtype: int64
```

```
df[ "Transmission" ] = le.fit_transform(df[ "Transmission" ])
df[ "Transmission" ].value_counts()
```

```
Transmission
```

```
15    1324
17    1211
```

```
25    901
3     789
5     490
9     445
6     339
16    319
19    295
24    193
12    168
8     132
22    118
21    113
26    91
2     84
18    77
1     65
10    62
4     53
23    39
0     31
14    26
20    11
7     4
11    3
13    2
Name: count, dtype: int64
```

```
df['Transmission'].value_counts()
```

```
Transmission
```

```
15    1324
17    1211
25    901
3     789
5     490
9     445
6     339
16    319
19    295
24    193
12    168
8     132
22    118
21    113
26    91
2     84
18    77
1     65
10    62
4     53
```

```

23      39
0       31
14      26
20      11
7        4
11      3
13      2
Name: count, dtype: int64

df['Fuel Type'].value_counts()

Fuel Type
X      3637
Z      3202
E      370
D      175
N       1
Name: count, dtype: int64

df["Fuel Type"] = le.fit_transform(df["Fuel Type"])
df["Fuel Type"].value_counts()

Fuel Type
3      3637
4      3202
1      370
0      175
2       1
Name: count, dtype: int64

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   Make             7385 non-null    int64  
 1   Vehicle Class   7385 non-null    int64  
 2   Engine Size(L)  7385 non-null    float64 
 3   Cylinders       7385 non-null    int64  
 4   Transmission    7385 non-null    int64  
 5   Fuel Type       7385 non-null    int64  
 6   Fuel Consumption City (L/100 km) 7385 non-null    float64 
 7   Fuel Consumption Hwy (L/100 km)   7385 non-null    float64 
 8   Fuel Consumption Comb (L/100 km)  7385 non-null    float64 
 9   Fuel Consumption Comb (mpg)       7385 non-null    int64  
 10  CO2 Emissions(g/km)            7385 non-null    int64  
dtypes: float64(4), int64(7)
memory usage: 634.8 KB

```

```

# df = df.drop(['Make', 'Model', 'Vehicle Class', 'Transmission'],
axis=1)
# df.shape

df["Fuel Type"].value_counts()

Fuel Type
3    3637
4    3202
1    370
0    175
2     1
Name: count, dtype: int64

from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df["Fuel Type"] = le.fit_transform(df["Fuel Type"])
df["Fuel Type"].value_counts()

Fuel Type
3    3637
4    3202
1    370
0    175
2     1
Name: count, dtype: int64

```

Correlation

```

correlation = df.corr()
correlation

          Make  Vehicle Class  Engine
Size(L) \
Make           1.000000      -0.029558      -
0.146199
Vehicle Class        -0.029558       1.000000
0.142704
Engine Size(L)        -0.146199       0.142704
1.000000
Cylinders           -0.162065       0.105978
0.927653
Transmission         0.181923      -0.156562      -
0.322389
Fuel Type            0.045368      -0.033560
0.058296
Fuel Consumption City (L/100 km) -0.197389       0.240941

```

| | | | |
|----------------------------------|------------------------------------|--------------|-----------|
| 0.831379 | | | |
| Fuel Consumption Hwy (L/100 km) | -0.126010 | 0.329828 | |
| 0.761526 | | | |
| Fuel Consumption Comb (L/100 km) | -0.175238 | 0.274388 | |
| 0.817060 | | | |
| Fuel Consumption Comb (mpg) | 0.182649 | -0.277606 | - |
| 0.757854 | | | |
| CO2 Emissions(g/km) | -0.151955 | 0.286468 | |
| 0.851145 | | | |
| | Cylinders | Transmission | Fuel |
| Type \ Make | -0.162065 | 0.181923 | 0.045368 |
| Vehicle Class | 0.105978 | -0.156562 | -0.033560 |
| Engine Size(L) | 0.927653 | -0.322389 | 0.058296 |
| Cylinders | 1.000000 | -0.270011 | 0.125175 |
| Transmission | -0.270011 | 1.000000 | 0.212872 |
| Fuel Type | 0.125175 | 0.212872 | 1.000000 |
| Fuel Consumption City (L/100 km) | 0.800702 | -0.345839 | -0.075605 |
| Fuel Consumption Hwy (L/100 km) | 0.715252 | -0.355371 | -0.129812 |
| Fuel Consumption Comb (L/100 km) | 0.780534 | -0.353609 | -0.095539 |
| Fuel Consumption Comb (mpg) | -0.719321 | 0.331213 | -0.016880 |
| CO2 Emissions(g/km) | 0.832644 | -0.316660 | 0.100306 |
| | Fuel Consumption City (L/100 km) \ | | |
| Make | | -0.197389 | |
| Vehicle Class | | 0.240941 | |
| Engine Size(L) | | 0.831379 | |
| Cylinders | | 0.800702 | |
| Transmission | | -0.345839 | |
| Fuel Type | | -0.075605 | |
| Fuel Consumption City (L/100 km) | | 1.000000 | |
| Fuel Consumption Hwy (L/100 km) | | 0.948180 | |
| Fuel Consumption Comb (L/100 km) | | 0.993810 | |
| Fuel Consumption Comb (mpg) | | -0.927059 | |
| CO2 Emissions(g/km) | | 0.919592 | |
| | Fuel Consumption Hwy (L/100 km) \ | | |
| Make | | -0.126010 | |
| Vehicle Class | | 0.329828 | |

| | |
|----------------------------------|-----------|
| Engine Size(L) | 0.761526 |
| Cylinders | 0.715252 |
| Transmission | -0.355371 |
| Fuel Type | -0.129812 |
| Fuel Consumption City (L/100 km) | 0.948180 |
| Fuel Consumption Hwy (L/100 km) | 1.000000 |
| Fuel Consumption Comb (L/100 km) | 0.977299 |
| Fuel Consumption Comb (mpg) | -0.890638 |
| CO2 Emissions(g/km) | 0.883536 |

| | Fuel Consumption Comb (L/100 km) \ |
|----------------------------------|------------------------------------|
| Make | -0.175238 |
| Vehicle Class | 0.274388 |
| Engine Size(L) | 0.817060 |
| Cylinders | 0.780534 |
| Transmission | -0.353609 |
| Fuel Type | -0.095539 |
| Fuel Consumption City (L/100 km) | 0.993810 |
| Fuel Consumption Hwy (L/100 km) | 0.977299 |
| Fuel Consumption Comb (L/100 km) | 1.000000 |
| Fuel Consumption Comb (mpg) | -0.925576 |
| CO2 Emissions(g/km) | 0.918052 |

| | Fuel Consumption Comb (mpg) \ |
|----------------------------------|-------------------------------|
| Make | 0.182649 |
| Vehicle Class | -0.277606 |
| Engine Size(L) | -0.757854 |
| Cylinders | -0.719321 |
| Transmission | 0.331213 |
| Fuel Type | -0.016880 |
| Fuel Consumption City (L/100 km) | -0.927059 |
| Fuel Consumption Hwy (L/100 km) | -0.890638 |
| Fuel Consumption Comb (L/100 km) | -0.925576 |
| Fuel Consumption Comb (mpg) | 1.000000 |
| CO2 Emissions(g/km) | -0.907426 |


| | CO2 Emissions(g/km) |
|----------------------------------|---------------------|
| Make | -0.151955 |
| Vehicle Class | 0.286468 |
| Engine Size(L) | 0.851145 |
| Cylinders | 0.832644 |
| Transmission | -0.316660 |
| Fuel Type | 0.100306 |
| Fuel Consumption City (L/100 km) | 0.919592 |
| Fuel Consumption Hwy (L/100 km) | 0.883536 |
| Fuel Consumption Comb (L/100 km) | 0.918052 |
| Fuel Consumption Comb (mpg) | -0.907426 |
| CO2 Emissions(g/km) | 1.000000 |

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Make              7385 non-null    int64  
 1   Vehicle Class     7385 non-null    int64  
 2   Engine Size(L)   7385 non-null    float64 
 3   Cylinders         7385 non-null    int64  
 4   Transmission      7385 non-null    int64  
 5   Fuel Type          7385 non-null    int64  
 6   Fuel Consumption City (L/100 km) 7385 non-null    float64 
 7   Fuel Consumption Hwy (L/100 km) 7385 non-null    float64  
 8   Fuel Consumption Comb (L/100 km) 7385 non-null    float64 
 9   Fuel Consumption Comb (mpg)       7385 non-null    int64  
 10  CO2 Emissions(g/km)        7385 non-null    int64  
dtypes: float64(4), int64(7)
memory usage: 634.8 KB
```

```
# constructing a heatmap to understand the correlation
plt.figure(figsize=(8,8))
sns.heatmap(correlation, square=True, fmt='.1f', annot=True)
```

```
<Axes: >
```



DATA SPLITTING

```
X = df.drop('CO2 Emissions(g/km)', axis=1)
Y = df['CO2 Emissions(g/km)']
```

```
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Make            7385 non-null   int64
```

```

1  Vehicle Class                7385 non-null    int64
2  Engine Size(L)              7385 non-null    float64
3  Cylinders                   7385 non-null    int64
4  Transmission                 7385 non-null    int64
5  Fuel Type                    7385 non-null    int64
6  Fuel Consumption City (L/100 km) 7385 non-null    float64
7  Fuel Consumption Hwy (L/100 km) 7385 non-null    float64
8  Fuel Consumption Comb (L/100 km) 7385 non-null    float64
9  Fuel Consumption Comb (mpg)     7385 non-null    int64
dtypes: float64(4), int64(6)
memory usage: 577.1 KB

Y.info()

<class 'pandas.core.series.Series'>
RangeIndex: 7385 entries, 0 to 7384
Series name: CO2 Emissions(g/km)
Non-Null Count Dtype
-----
7385 non-null    int64
dtypes: int64(1)
memory usage: 57.8 KB

from sklearn.model_selection import train_test_split
X_TRAIN , X_TEST , Y_TRAIN, Y_TEST = train_test_split(X,Y, test_size = 0.25, random_state=25)
print("Size of Train X = " , len(X_TRAIN))
print("Size of Train Y = " , len(Y_TRAIN))
print("Size of Test X = " , len(X_TEST))
print("Size of Test Y = " , len(Y_TEST))

Size of Train X =  5538
Size of Train Y =  5538
Size of Test X =  1847
Size of Test Y =  1847

```

LINEAR REGRESSION

```

from sklearn.linear_model import LinearRegression
model= LinearRegression()
model.fit(X_TRAIN, Y_TRAIN)

LinearRegression()

```

Prediction on Train Data

```
# accuracy for prediction on training data
training_data_prediction = model.predict(X_TRAIN)
print(training_data_prediction)

[280.013408  315.84268011 144.72620894 ... 294.17432435 236.5539138
 198.74325804]

# R squared error
score_1 = metrics.r2_score(Y_TRAIN, training_data_prediction)

# Mean Absolute Error
score_2 = metrics.mean_absolute_error(Y_TRAIN,
training_data_prediction)

print("R squared : ", score_1)
print('Mean Absolute Error : ', score_2)

R squared :  0.9153849568066784
Mean Absolute Error :  11.17147303195528

plt.scatter(Y_TRAIN, training_data_prediction)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual Price vs Preicted Price")
plt.show()
```



Prediction on Test Data

```

y_pred = model.predict(X_TEST)
y_pred
print(y_pred)

[233.2941932 283.77211363 242.956709 ... 237.3456667 191.57369611
 174.13609452]

# R squared Score
score_1 = metrics.r2_score(Y_TEST, y_pred)

# Mean Absolute Error
score_2 = metrics.mean_absolute_error(Y_TEST, y_pred)


print("R squared Score : ", score_1)
print('Mean Absolute Error : ', score_2)

R squared Score :  0.9173989134759174
Mean Absolute Error :  11.044093610214464

plt.scatter(Y_TEST, y_pred)
plt.xlabel("Actual")

```

```
plt.ylabel("Predicted")
plt.title("Actual vs Preicted")
plt.show()
```



SVM REGRESSION

```
from sklearn.svm import SVR
svr = SVR(kernel='linear')

svr.fit(X_TRAIN, Y_TRAIN)


SVR(kernel='linear')
```

Prediction on Train Data

```
# accuracy for prediction on training data
training_data_prediction = svr.predict(X_TRAIN)
print(training_data_prediction)
```

```
[279.58450148 319.15835613 150.61361469 ... 290.65433483 279.41839813  
193.94021626]
```

```
# R squared error  
score_1 = metrics.r2_score(Y_TRAIN, training_data_prediction)  
  
# Mean Absolute Error  
score_2 = metrics.mean_absolute_error(Y_TRAIN,  
training_data_prediction)  
  
print("R squared : ", score_1)  
print('Mean Absolute Error : ', score_2)  
  
R squared : 0.812235212106675  
Mean Absolute Error : 8.393098579369461  
  
plt.scatter(Y_TRAIN, training_data_prediction)  
plt.xlabel("Actual Prices")  
plt.ylabel("Predicted Prices")  
plt.title("Actual Price vs Preicted Price")  
plt.show()
```



Prediction on Test Data

```
y_pred = svr.predict(X_TEST)
y_pred
print(y_pred)

[241.93941487 283.5429716 238.16100842 ... 233.67655131 185.32175089
 174.78062399]


# R squared Score
score_1 = metrics.r2_score(Y_TEST, y_pred)

# Mean Absolute Error
score_2 = metrics.mean_absolute_error(Y_TEST, y_pred)

print("R squared Score : ", score_1)
print('Mean Absolute Error : ', score_2)

R squared Score :  0.819230911133283
Mean Absolute Error :  8.066009439876659

plt.scatter(Y_TEST, y_pred)
plt.xlabel("Actual")
plt.ylabel("Predicted")
plt.title("Actual vs Predicted")
plt.show()
```



REGRESSION DECISION TREE

```
# import the regressor
from sklearn.tree import DecisionTreeRegressor

# create a regressor object
regressor = DecisionTreeRegressor(random_state = 25)

# fit the regressor with X and Y data
regressor.fit(X_TRAIN, Y_TRAIN)

DecisionTreeRegressor(random_state=25)
```

Prediction on Train Data

```
# accuracy for prediction on training data
training_data_prediction = regressor.predict(X_TRAIN)
print(training_data_prediction)

[279.          321.          150.66666667 ... 290.          195.]
```

```

# R squared error
score_1 = metrics.r2_score(Y_TRAIN, training_data_prediction)


# Mean Absolute Error
score_2 = metrics.mean_absolute_error(Y_TRAIN,
training_data_prediction)

print("R squared : ", score_1)
print('Mean Absolute Error : ', score_2)

R squared :  0.9996932983241511
Mean Absolute Error :  0.37144834820891154

plt.scatter(Y_TRAIN, training_data_prediction)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual Price vs Preicted Price")
plt.show()

```



Prediction on Test Data

```
y_pred = regressor.predict(X_TEST)
y_pred
print(y_pred)

[244. 285. 238. ... 233. 186. 172.5]


# R squared Score
score_1 = metrics.r2_score(Y_TEST, y_pred)

# Mean Absolute Error
score_2 = metrics.mean_absolute_error(Y_TEST, y_pred)

print("R squared Score : ", score_1)
print('Mean Absolute Error : ', score_2)

R squared Score : 0.9975867297567413
Mean Absolute Error : 1.6402377085105833

plt.scatter(Y_TEST, y_pred)
plt.xlabel("Actual")
plt.ylabel("Predicted")
plt.title("Actual vs Predicted")
plt.show()
```



RANDOM FOREST

```
from sklearn.ensemble import RandomForestRegressor  
regressor = RandomForestRegressor(n_estimators=5, random_state=25)  
# fit the regressor with X and Y data  
regressor.fit(X_TRAIN, Y_TRAIN)  
RandomForestRegressor(n_estimators=5, random_state=25)
```

Prediction on Training Data

```
# accuracy for prediction on training data  
training_data_prediction = regressor.predict(X_TRAIN)  
print(training_data_prediction)
```

```

[278.2      321.      151.65333333 ... 290.      182.4
 195.      ]
# R squared error
score_1 = metrics.r2_score(Y_TRAIN, training_data_prediction)


# Mean Absolute Error
score_2 = metrics.mean_absolute_error(Y_TRAIN,
training_data_prediction)

print("R squared error : ", score_1)
print('Mean Absolute Error : ', score_2)

R squared error :  0.9991172692628274
Mean Absolute Error :  0.8785836894937649

plt.scatter(Y_TRAIN, training_data_prediction)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual Price vs Preicted Price")
plt.show()

```



Prediction on Test Data

```
y_pred = regressor.predict(X_TEST)
y_pred
array([244.          , 285.          , 238.          , ...,
       186.          , 174.06666667])


# R squared Score
score_1 = metrics.r2_score(Y_TEST, y_pred)

# Mean Absolute Error
score_2 = metrics.mean_absolute_error(Y_TEST, y_pred)

print("R squared Score : ", score_1)
print('Mean Absolute Error : ', score_2)

R squared Score :  0.9975115254335551
Mean Absolute Error :  1.818620193498374


plt.scatter(Y_TEST, y_pred)
plt.xlabel("Actual")
plt.ylabel("Predicted")
plt.title("Actual vs Predicted")
plt.show()
```



```
from sklearn.tree import plot_tree
import matplotlib.pyplot as plt

# Assuming regressor is your trained Random Forest model
# Pick one tree from the forest, e.g., the first tree (index 0)
tree_to_plot = regressor.estimators_[0]

# Plot the decision tree
plt.figure(figsize=(20, 10))
plot_tree(tree_to_plot, feature_names=df.columns.tolist(),
          filled=True, rounded=True, fontsize=10)
plt.title("Decision Tree from Random Forest")
plt.show()
```



xgboost

```
import xgboost as xgb
```

```
ModuleNotFoundError                         Traceback (most recent call
last)
Cell In[114], line 1
----> 1 import xgboost as xgb
```

ModuleNotFoundError: No module named 'xgboost'

```
xgb_model = xgb.XGBRegressor(objective='reg:squarederror', #  
Specify the objective for regression  
    n_estimators=100,  
    learning_rate=0.1,  
    max_depth=3,  
    random_state=42)
```

```
# Train the model  
xgb_model.fit(X_TRAIN, Y_TRAIN)
```

```
4                         max_depth=3,  
5                         random_state=42)  
7 # Train the model  
8 xgb_model.fit(X_TRAIN, Y_TRAIN)
```

NameError: name 'xgb' is not defined